

The Linux Serial Programming HOWTO
by Peter H. Baumann, Peter.Baumann@dlr.de
v1.0, 22 January 1998

This document describes how to program communications with devices
over a serial port on a Linux box.

Table of Contents

1. Introduction

- 1.1 Copyright
- 1.2 New Versions Of This Document
- 1.3 Feedback

2. Getting started

- 2.1 Debugging
- 2.2 Port Settings
- 2.3 Input Concepts for Serial Devices
 - 2.3.1 Canonical Input Processing
 - 2.3.2 Non-Canonical Input Processing
 - 2.3.3 Asynchronous Input
 - 2.3.4 Waiting for Input from Multiple Sources

3. Program Examples

- 3.1 Canonical Input Processing
- 3.2 Non-Canonical Input Processing
- 3.3 Asynchronous Input
- 3.4 Waiting for Input from Multiple Sources

4. Other Sources of Information

5. Contributions

1. Introduction

This is the Linux Serial Programming HOWTO. All about how to program communications with other devices / computers over a serial line under Linux. Different techniques are explained: Canonical I/O (only complete lines are transmitted/received), asynchronous I/O, and waiting for input from multiple sources.

This document does not describe how to set up serial ports, because this has been described by Greg Hankins in the Serial-HOWTO.

I have to emphasize that I am not an expert in this field, but have had problems with a project that involved such communication. The code examples presented here were derived from the miniterm code available from the LDP programmers guide (<ftp://sunsite.unc.edu/pub/Linux/docs/LDP/programmers-guide/lpg-0.4.tar.gz> and mirrors) in the examples directory.

Since I wrote this document in June 1997, I have moved to WinNT to satisfy customers need, so I have not built up more in depth

knowledge. If anybody has any comments, I will gladly incorporate them into this document (see sect. Feedback). If someone would like to take over and do a better job, please e-mail me.

All examples were tested using a i386 Linux Kernel 2.0.29.

1.1. Copyright

The Linux Serial-Programming-HOWTO is copyright (C) 1997 by Peter Baumann. Linux HOWTO documents may be reproduced and distributed in whole or in part, in any medium physical or electronic, as long as this copyright notice is retained on all copies. Commercial redistribution is allowed and encouraged; however, the author would like to be notified of any such distributions.

All translations, derivative works, or aggregate works incorporating any Linux HOWTO documents must be covered under this copyright notice. That is, you may not produce a derivative work from a HOWTO and impose additional restrictions on its distribution. Exceptions to these rules may be granted under certain conditions; please contact the Linux HOWTO coordinator at the address given below.

In short, we wish to promote dissemination of this information through as many channels as possible. However, we do wish to retain copyright on the HOWTO documents, and would like to be notified of any plans to redistribute the HOWTOs.

If you have questions, please contact Tim Bynum, the Linux HOWTO coordinator, at linux-howto@sunsite.unc.edu via email.

1.2. New Versions Of This Document

New versions of the Serial-Programming-HOWTO will be available at and mirror sites. There are other formats, such as PostScript and DVI versions in the other-formats directory. The Serial-Programming-HOWTO is also available at <http://sunsite.unc.edu/LDP/HOWTO/Serial-Programming-HOWTO.html> and will be posted to comp.os.linux.answers monthly.

1.3. Feedback

Please send me any corrections, questions, comments, suggestions, or additional material. I would like to improve this HOWTO! Tell me exactly what you don't understand, or what could be clearer. You can reach me at Peter.Baumann@dlr.de via email. Please include the version number of the Serial-Programming-HOWTO when writing, this is version 0.3.

2. Getting started

2.1. Debugging

The best way to debug your code is to set up another Linux box, and connect the two computers via a null-modem cable. Use miniterm (available from the LDP programmers guide (<ftp://sunsite.unc.edu/pub/Linux/docs/LDP/programmers-guide/lpg-0.4.tar.gz> in the examples directory) to transmit characters to your Linux box. Miniterm can be compiled very easily and will transmit all keyboard input raw over the serial port. Only the define statement `#define MODEMDEVICE "/dev/ttyS0"` has to be checked. Set it

to ttyS0 for COM1, ttyS1 for COM2, etc.. It is essential for testing, that all characters are transmitted raw (without output processing) over the line. To test your connection, start miniterm on both computers and just type away. The characters input on one computer should appear on the other computer and vice versa. The input will not be echoed to the attached screen.

To make a null-modem cable you have to cross the TxD (transmit) and RxD (receive) lines. For a description of a cable see sect. 7 of the Serial-HOWTO.

It is also possible to perform this testing with only one computer, if you have two unused serial ports. You can then run two miniterms off two virtual consoles. If you free a serial port by disconnecting the mouse, remember to redirect /dev/mouse if it exists. If you use a multiport serial card, be sure to configure it correctly. I had mine configured wrong and everything worked fine as long as I was testing only on my computer. When I connected to another computer, the port started loosing characters. Executing two programs on one computer just isn't fully asynchronous.

2.2. Port Settings

The devices /dev/ttyS* are intended to hook up terminals to your Linux box, and are configured for this use after startup. This has to be kept in mind when programming communication with a raw device. E.g. the ports are configured to echo characters sent from the device back to it, which normally has to be changed for data transmission.

All parameters can be easily configured from within a program. The configuration is stored in a structure struct termios, which is defined in <asm/termbits.h>:

```
#define NCCS 19
struct termios {
    tcflag_t c_iflag;           /* input mode flags */
    tcflag_t c_oflag;           /* output mode flags */
    tcflag_t c_cflag;           /* control mode flags */
    tcflag_t c_lflag;           /* local mode flags */
    cc_t c_line;                /* line discipline */
    cc_t c_cc[NCCS];            /* control characters */
};
```

This file also includes all flag definitions. The input mode flags in c_iflag handle all input processing, which means that the characters sent from the device can be processed before they are read with read. Similarly c_oflag handles the output processing. c_cflag contains the settings for the port, as the baudrate, bits per character, stop bits, etc.. The local mode flags stored in c_lflag determine if characters are echoed, signals are sent to your program, etc.. Finally the array c_cc defines the control characters for end of file, stop, etc.. Default values for the control characters are defined in <asm/termios.h>. The flags are described in the manual page termios(3). The structure termios contains the c_line (line discipline) element, which is not used in POSIX compliant systems.

2.3. Input Concepts for Serial Devices

Here three different input concepts will be presented. The appropriate concept has to be chosen for the intended application. Whenever

possible, do not loop reading single characters to get a complete string. When I did this, I lost characters, whereas a read for the whole string did not show any errors.

2.3.1. Canonical Input Processing

This is the normal processing mode for terminals, but can also be useful for communicating with other devices. Input is processed in units of lines, which means that a read will only return a full line of input. A line is by default terminated by a NL (ASCII LF), an end of file, or an end of line character. A CR (the DOS/Windows default end-of-line) will not terminate a line with the default settings.

Canonical input processing can also handle the erase, delete word, and reprint characters, translate CR to NL, etc..

2.3.2. Non-Canonical Input Processing

Non-Canonical Input Processing will handle a fixed amount of characters per read, and allows for a character timer. This mode should be used if your application will always read a fixed number of characters, or if the connected device sends bursts of characters.

2.3.3. Asynchronous Input

The two modes described above can be used in synchronous and asynchronous mode. Synchronous is the default, where a read statement will block, until the read is satisfied. In asynchronous mode the read statement will return immediately and send a signal to the calling program upon completion. This signal can be received by a signal handler.

2.3.4. Waiting for Input from Multiple Sources

This is not a different input mode, but might be useful, if you are handling multiple devices. In my application I was handling input over a TCP/IP socket and input over a serial connection from another computer quasi-simultaneously. The program example given below will wait for input from two different input sources. If input from one source becomes available, it will be processed, and the program will then wait for new input.

The approach presented below seems rather complex, but it is important to keep in mind that Linux is a multi-processing operating system. The select system call will not load the CPU while waiting for input, whereas looping until input becomes available would slow down other processes executing at the same time.

3. Program Examples

All examples have been derived from miniterm.c. The type ahead buffer is limited to 255 characters, just like the maximum string length for canonical input processing (`<linux/limits.h>` or `<posix1_lim.h>`).

See the comments in the code for explanation of the use of the different input modes. I hope that the code is understandable. The example for canonical input is commented best, the other examples are commented only where they differ from the example for canonical input to emphasize the differences.

The descriptions are not complete, but you are encouraged to

experiment with the examples to derive the best solution for your application.

Don't forget to give the appropriate serial ports the right permissions (e. g.: `chmod a+rw /dev/ttyS1`)!

3.1. Canonical Input Processing

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <termios.h>
#include <stdio.h>

/* baudrate settings are defined in <asm/termbits.h>, which is
included by <termios.h> */
#define BAUDRATE B38400
/* change this definition for the correct port */
#define MODEMDEVICE "/dev/ttyS1"
#define _POSIX_SOURCE 1 /* POSIX compliant source */

#define FALSE 0
#define TRUE 1

volatile int STOP=FALSE;

main()
{
    int fd,c, res;
    struct termios oldtio,newtio;
    char buf[255];
    /*
    Open modem device for reading and writing and not as controlling tty
    because we don't want to get killed if linenoise sends CTRL-C.
    */
    fd = open(MODEMDEVICE, O_RDWR | O_NOCTTY );
    if (fd <0) {perror(MODEMDEVICE); exit(-1); }

    tcgetattr(fd,&oldtio); /* save current serial port settings */
    bzero(&newtio, sizeof(newtio)); /* clear struct for new port settings */

    /*
    BAUDRATE: Set bps rate. You could also use cfsetispeed and cfsetospeed.
    CRTSCTS : output hardware flow control (only used if the cable has
               all necessary lines. See sect. 7 of Serial-HOWTO)
    CS8      : 8n1 (8bit,no parity,1 stopbit)
    CLOCAL   : local connection, no modem control
    CREAD    : enable receiving characters
    */
    newtio.c_cflag = BAUDRATE | CRTSCTS | CS8 | CLOCAL | CREAD;

    /*
    IGNPAR   : ignore bytes with parity errors
    ICRNL    : map CR to NL (otherwise a CR input on the other computer
               will not terminate input)
    otherwise make device raw (no other input processing)
    */
    newtio.c_iflag = IGNPAR | ICRNL;

    /*
```

```

Raw output.
*/
newtio.c_oflag = 0;

/*
  ICANON : enable canonical input
  disable all echo functionality, and don't send signals to calling program
*/
newtio.c_lflag = ICANON;

/*
  initialize all control characters
  default values can be found in /usr/include/termios.h, and are given
  in the comments, but we don't need them here
*/
newtio.c_cc[VINTR]    = 0;    /* Ctrl-c */
newtio.c_cc[VQUIT]    = 0;    /* Ctrl-\ */
newtio.c_cc[VERASE]    = 0;    /* del */
newtio.c_cc[VKILL]    = 0;    /* @ */
newtio.c_cc[VEOF]     = 4;    /* Ctrl-d */
newtio.c_cc[VTIME]    = 0;    /* inter-character timer unused */
newtio.c_cc[VMIN]     = 1;    /* blocking read until 1 character arrives */
newtio.c_cc[VSWTC]    = 0;    /* '\0' */
newtio.c_cc[VSTART]   = 0;    /* Ctrl-q */
newtio.c_cc[VSTOP]    = 0;    /* Ctrl-s */
newtio.c_cc[VSUSP]    = 0;    /* Ctrl-z */
newtio.c_cc[VEOL]     = 0;    /* '\0' */
newtio.c_cc[VREPRINT] = 0;    /* Ctrl-r */
newtio.c_cc[VDISCARD] = 0;    /* Ctrl-u */
newtio.c_cc[VWERASE]  = 0;    /* Ctrl-w */
newtio.c_cc[VLNEXT]   = 0;    /* Ctrl-v */
newtio.c_cc[VEOL2]    = 0;    /* '\0' */

/*
  now clean the modem line and activate the settings for the port
*/
tcflush(fd, TCIFLUSH);
tcsetattr(fd, TCSANOW, &newtio);

/*
  terminal settings done, now handle input
  In this example, inputting a 'z' at the beginning of a line will
  exit the program.
*/
while (STOP==FALSE) {    /* loop until we have a terminating condition */
/* read blocks program execution until a line terminating character is
input, even if more than 255 chars are input. If the number
of characters read is smaller than the number of chars available,
subsequent reads will return the remaining chars. res will be set
to the actual number of characters actually read */
res = read(fd, buf, 255);
buf[res]=0;    /* set end of string, so we can printf */
printf(":%s:%d\n", buf, res);
if (buf[0]=='z') STOP=TRUE;
}
/* restore the old port settings */
tcsetattr(fd, TCSANOW, &oldtio);
}

```

3.2. Non-Canonical Input Processing

In non-canonical input processing mode, input is not assembled into lines and input processing (erase, kill, delete, etc.) does not occur. Two parameters control the behavior of this mode: `c_cc[VTIME]` sets the character timer, and `c_cc[VMIN]` sets the minimum number of characters to receive before satisfying the read.

If `MIN > 0` and `TIME = 0`, `MIN` sets the number of characters to receive before the read is satisfied. As `TIME` is zero, the timer is not used.

If `MIN = 0` and `TIME > 0`, `TIME` serves as a timeout value. The read will be satisfied if a single character is read, or `TIME` is exceeded ($t = \text{TIME} * 0.1 \text{ s}$). If `TIME` is exceeded, no character will be returned.

If `MIN > 0` and `TIME > 0`, `TIME` serves as an inter-character timer. The read will be satisfied if `MIN` characters are received, or the time between two characters exceeds `TIME`. The timer is restarted every time a character is received and only becomes active after the first character has been received.

If `MIN = 0` and `TIME = 0`, read will be satisfied immediately. The number of characters currently available, or the number of characters requested will be returned. According to Antonino (see contributions), you could issue a `fcntl(fd, F_SETFL, FNDELAY)`; before reading to get the same result.

By modifying `newtio.c_cc[VTIME]` and `newtio.c_cc[VMIN]` all modes described above can be tested.

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <termios.h>
#include <stdio.h>

#define BAUDRATE B38400
#define MODEMDEVICE "/dev/ttyS1"
#define _POSIX_SOURCE 1 /* POSIX compliant source */
#define FALSE 0
#define TRUE 1

volatile int STOP=FALSE;

main()
{
    int fd,c, res;
    struct termios oldtio,newtio;
    char buf[255];

    fd = open(MODEMDEVICE, O_RDWR | O_NOCTTY );
    if (fd <0) {perror(MODEMDEVICE); exit(-1); }

    tcgetattr(fd,&oldtio); /* save current port settings */

    bzero(&newtio, sizeof(newtio));
    newtio.c_cflag = BAUDRATE | CRTSCTS | CS8 | CLOCAL | CREAD;
    newtio.c_iflag = IGNPAR;
    newtio.c_oflag = 0;

    /* set input mode (non-canonical, no echo,...) */
```

```

newtio.c_lflag = 0;

newtio.c_cc[VTIME]    = 0;    /* inter-character timer unused */
newtio.c_cc[VMIN]     = 5;    /* blocking read until 5 chars received */

tcflush(fd, TCIFLUSH);
tcsetattr(fd,TCSANOW,&newtio);

while (STOP==FALSE) {      /* loop for input */
    res = read(fd,buf,255); /* returns after 5 chars have been input */
    buf[res]=0;             /* so we can printf... */
    printf(":%s:%d\n", buf, res);
    if (buf[0]=='z') STOP=TRUE;
}
tcsetattr(fd,TCSANOW,&oldtio);
}

```

3.3. Asynchronous Input

```

#include <termios.h>
#include <stdio.h>
#include <unistd.h>
#include <fcntl.h>
#include <sys/signal.h>
#include <sys/types.h>

#define BAUDRATE B38400
#define MODEMDEVICE "/dev/ttyS1"
#define _POSIX_SOURCE 1 /* POSIX compliant source */
#define FALSE 0
#define TRUE 1

volatile int STOP=FALSE;

void signal_handler_IO (int status); /* definition of signal handler */
int wait_flag=TRUE;                /* TRUE while no signal received */

main()
{
    int fd,c, res;
    struct termios oldtio,newtio;
    struct sigaction saio;          /* definition of signal action */
    char buf[255];

    /* open the device to be non-blocking (read will return immediatly) */
    fd = open(MODEMDEVICE, O_RDWR | O_NOCTTY | O_NONBLOCK);
    if (fd <0) {perror(MODEMDEVICE); exit(-1); }

    /* install the signal handler before making the device asynchronous */
    saio.sa_handler = signal_handler_IO;
    saio.sa_mask = 0;
    saio.sa_flags = 0;
    saio.sa_restorer = NULL;
    sigaction(SIGIO,&saio,NULL);

    /* allow the process to receive SIGIO */
    fcntl(fd, F_SETOWN, getpid());
    /* Make the file descriptor asynchronous (the manual page says only
       O_APPEND and O_NONBLOCK, will work with F_SETFL...) */
    fcntl(fd, F_SETFL, FASYNC);

```



```

tcgetattr(fd,&oldtio); /* save current port settings */
/* set new port settings for canonical input processing */
newtio.c_cflag = BAUDRATE | CRTSCTS | CS8 | CLOCAL | CREAD;
newtio.c_iflag = IGNPAR | ICRNL;
newtio.c_oflag = 0;
newtio.c_lflag = ICANON;
newtio.c_cc[VMIN]=1;
newtio.c_cc[VTIME]=0;
tcflush(fd, TCIFLUSH);
tcsetattr(fd,TCSANOW,&newtio);

/* loop while waiting for input. normally we would do something
   useful here */
while (STOP==FALSE) {
    printf(".\n");usleep(100000);
    /* after receiving SIGIO, wait_flag = FALSE, input is available
       and can be read */
    if (wait_flag==FALSE) {
        res = read(fd,buf,255);
        buf[res]=0;
        printf(":%s:%d\n", buf, res);
        if (res==1) STOP=TRUE; /* stop loop if only a CR was input */
        wait_flag = TRUE;      /* wait for new input */
    }
}
/* restore old port settings */
tcsetattr(fd,TCSANOW,&oldtio);
}

/*****
* signal handler. sets wait_flag to FALSE, to indicate above loop that
* characters have been received.
*****/

void signal_handler_IO (int status)
{
    printf("received SIGIO signal.\n");
    wait_flag = FALSE;
}

```

3.4. Waiting for Input from Multiple Sources

This section is kept to a minimum. It is just intended to be a hint, and therefore the example code is kept short. This will not only work with serial ports, but with any set of file descriptors.

The select call and accompanying macros use a fd_set. This is a bit array, which has a bit entry for every valid file descriptor number. select will accept a fd_set with the bits set for the relevant file descriptors and returns a fd_set, in which the bits for the file descriptors are set where input, output, or an exception occurred. All handling of fd_set is done with the provided macros. See also the manual page select(2).

```

#include <sys/time.h>
#include <sys/types.h>
#include <unistd.h>

main()

```

```

{
    int    fd1, fd2; /* input sources 1 and 2 */
    fd_set readfs; /* file descriptor set */
    int    maxfd; /* maximum file descriptor used */
    int    loop=1; /* loop while TRUE */

    /* open_input_source opens a device, sets the port correctly, and
       returns a file descriptor */
    fd1 = open_input_source("/dev/ttyS1"); /* COM2 */
    if (fd1<0) exit(0);
    fd2 = open_input_source("/dev/ttyS2"); /* COM3 */
    if (fd2<0) exit(0);
    maxfd = MAX (fd1, fd2)+1; /* maximum bit entry (fd) to test */

    /* loop for input */
    while (loop) {
        FD_SET(fd1, &readfs); /* set testing for source 1 */
        FD_SET(fd2, &readfs); /* set testing for source 2 */
        /* block until input becomes available */
        select(maxfd, &readfs, NULL, NULL, NULL);
        if (FD_ISSET(fd1)) /* input from source 1 available */
            handle_input_from_source1();
        if (FD_ISSET(fd2)) /* input from source 2 available */
            handle_input_from_source2();
    }
}

```

The given example blocks indefinitely, until input from one of the sources becomes available. If you need to timeout on input, just replace the select call by:

```

    int res;
    struct timeval Timeout;

    /* set timeout value within input loop */
    Timeout.tv_usec = 0; /* milliseconds */
    Timeout.tv_sec = 1; /* seconds */
    res = select(maxfd, &readfs, NULL, NULL, &Timeout);
    if (res==0)
        /* number of file descriptors with input = 0, timeout occurred. */

```

This example will timeout after 1 second. If a timeout occurs, select will return 0, but beware that Timeout is decremented by the time actually waited for input by select. If the timeout value is zero, select will return immediately.

4. Other Sources of Information

- The Linux Serial-HOWTO describes how to set up serial ports and contains hardware information.
- Serial Programming Guide for POSIX Compliant Operating Systems <<http://www.easysw.com/~mike/serial>>, by Michael Sweet. This link is obsolete and I could not find a new location for it. Does somebody know where we can find it again? It was a well prepared document!
- The manual page termios(3) describes all flags for the termios structure.

5. Contributions

As mentioned in the introduction, I am no expert in this field, but had problems myself, and found a solution with the help of others. Thanks for the help from Mr. Strudthoff from the European Transonic Windtunnel, Cologne, Michael Carter (mcarter@rocke.electro.swri.edu), and Peter Waltenberg (p.waltenberg@karaka.chch.cri.nz)

Antonino Ianella (antonino@usa.net) wrote the Serial-Port-Programming Mini HOWTO, at the same time I prepared this document. Greg Hankins asked me to incorporate Antonino's Mini-HOWTO into this document.

The structure of this document and SGML formatting was derived from the Serial-HOWTO by Greg Hankins. Thanks also for various corrections made by : Dave Pfaltzgraff (Dave_Pfaltzgraff@patapsco.com), Sean Lincolne (slincol@tpgi.com.au), Michael Wiedmann (mw@miwie.in-berlin.de), and Adrey Bonar (andy@tipas.lt).